

Sequence-to-sequence models: Part II

Seq2Seq / encoder-decoder models

Ryan Chan

3 May 2023

Outline

Recap of Part I: RNNs and LSTMs

RNN variants: stacked & bidirectional

RNN/LSTM applications

Sequence-to-sequence (Seq2Seq) / Encoder-decoder RNNs
Seq2Seq architecture


Neural machine translation

Training Seq2Seq

Weaknesses of Seq2Seq

Recap: Recurrent Neural Networks

- RNNs [Hopfield, 1982; Rumelhart et al., 1985] are capable of conditioning the model on *all* previous tokens (in theory)




- Core idea: the hidden layer from previous timestep provides a form of **memory** or **context** that informs decisions to be made later in the sequence
- The **same** weights are applied at every timestep

¹Manning et al. [2017]

Recap: Training RNN language models

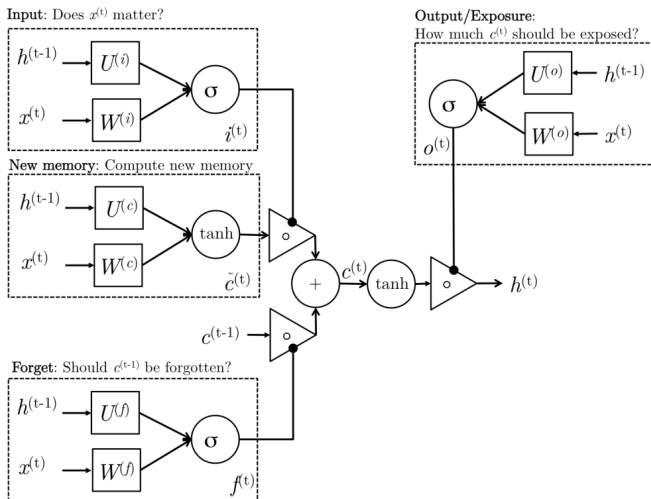
- We use the **cross-entropy** loss between the output distribution $\hat{\mathbf{y}}^{(t)}$ and $\mathbf{y}^{(t)}$ (and average over the entire batch/corpus)
- We use **teacher-forcing**: input the previous words/tokens to predict the next



Recap: LSTMs

- “Vanilla” RNNs can be **difficult to train**
 - Vanishing/exploding gradients mean its hard to preserve memory over many timesteps
 - The hidden layers/states and weights are asked to do a *lot* of work
- Long short-term memory (LSTM) networks [Hochreiter and Schmidhuber, 1997; Gers et al., 2000] are the most commonly used extension to RNNs
 - In practice, could preserve information to about 100 timesteps rather than 7 in “vanilla” RNNs
 - See “**Understanding LSTM Networks**” [Olah, 2015] for a bit more of a deeper explanation and walkthrough


Recap: LSTM architecture



3

Stacked / Deep RNNs


- We can use the entire sequence of outputs from one RNN as the input of another RNN to create a **stack** of RNNs



- Outputs of one layer serves as input to a subsequent layer
- Typically outperforms single-layer networks but increases training cost


Bidirectional RNNs

- Run **two** separate RNNs: left-to-right (we've seen above) and right-to-left
- Then **concatenate** the hidden states from both RNNs
 - Alternatives: element-wise mean or sum




- Effective if you have access to a full input sentence (to **encode** a sentence)
- “BERT”: **Bidirectional** Encoder Representations from Transformers


RNN tasks




a) sequence labeling



b) sequence classification




c) language modeling



d) encoder-decoder

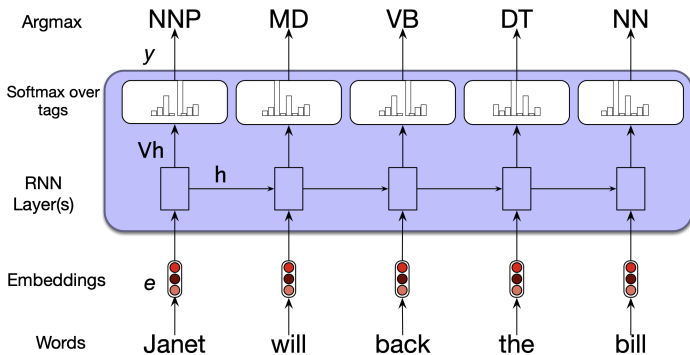
Language modelling

- **Causal language modelling / autoregressive generation**: incremental generation of words by repeated sampling of the next word conditional on previous choices
 - $\langle s \rangle$: beginning of sentence marker
 - $\langle /s \rangle$: end of sentence marker




Sequence labelling

- In **sequence labelling**, we aim to assign a label to each element in the sequence:
 - e.g. **part-of-speech tagging**, **named entity recognition**



Sequence classification


- In **sequence classification**, we use the RNN to classify the **entire sequence** (rather than the individual elements/tokens):
 - e.g. **sentiment analysis**, **spam detection**, **document-level topic classification**
- The loss comes from the final task, and we backpropagate through the entire network, i.e. **end-to-end training**



9


Sequence classification (with bidirectional RNN)

- If we have access to the full sequence, a **bidirectional RNN** can provide a more effective encoding of the sequence



Seq2Seq / Encoder-decoder RNN

- An architecture which takes in an input sequence and outputs a sequence that is of **different** length
 - Sequence labelling (part-of-speech tagging / named entity recognition) have two sequences that have the same length
- Some example use cases:
 - **Machine translation, summarisation, question answering**



Seq2Seq architecture

- Able to generate context appropriate, arbitrary length output sequences *given* an input sequence
 - **Encoder network**: takes in input sequence to produce a **contextualised representation** of it
 - **Decoder network**: takes in the context produced to generate a **text-specific output sequence**
- For example, an RNN Seq2Seq model:
 - The *encoder RNN* produces a **context**

$$\mathbf{c} = f(\mathbf{h}_{1:n}^e) \quad (1)$$

where $\mathbf{h}_{1:n}^e = \{\mathbf{h}_1^e, \dots, \mathbf{h}_n^e\}$ are the hidden states of the encoder RNN

- The *decoder RNN* accepts \mathbf{c} and generates a sequence of hidden states $\mathbf{h}_{1:m}^d$ to form corresponding sequence of output states $\mathbf{y}_{1:m}$ (where possibly $m \neq n$)

Machine translation

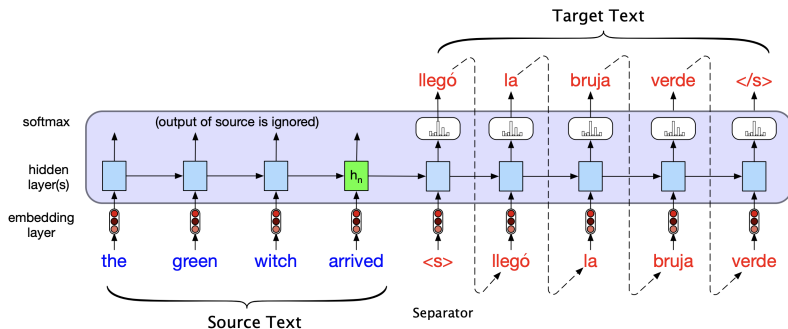
- **Machine translation (MT)** is the task of translating a sentence x in one language (**source language**) to a sentence y (**target language**)
- Early 1950s: MT research started but relied heavily on simple rule-based systems using word substitution
- 1990s-2010s: **Statistical machine translation (SMT)**
 - Want to find the “best” target sentence y given the source sentence x :

$$\operatorname{argmax}_y p(y|x) = \operatorname{argmax}_y p(x|y)p(y) \quad (2)$$

- $p(x|y)$: translation model whose goal is to model how phrases are translated (trained with sentence pairs)
 - $p(y)$: language model whose goal is to write sensible sentences
- Systems were extremely complex which required lots of feature engineering and human effort to maintain
- Repeated effort for different language pairs

Seq2Seq for machine translation

- **Neural machine translation (NMT)**: use a single end-to-end neural network
- Each training example is a pair of strings: source, target
 - Concatenate with a separator token, e.g. $\langle s \rangle$, $\langle \text{EOS} \rangle$, etc.



Seq2Seq for machine translation

- The sequence-to-sequence model for NMT is an example of a **conditional language model**
 - **Language model**: the decoder is predicting the next word in the target sentence y
 - **Conditional**: the predictions are also *conditioned* on some source sentence x
- NMT *directly* estimates $p(y|x)$:


$$p(y|x) = p(y_1|x) \cdot p(y_2|y_1, x) \cdot p(y_3|y_{1:2}, x) \cdots p(y_T|y_{1:T-1}, x) \quad (3)$$

Seq2Seq successes in machine translation

- **BLEU: Bilingual Evaluation Understudy**
 - Compares a machine-written translation to one or several human-written translation(s) and computes a **similarity score** based on n -gram precision
 - Adds a penalty for translations which are too short
- BLEU is useful but not perfect: there are many valid ways to translate a sentence

MT progress over time

[Edinburgh En-De WMT newstest2013 Cased BLEU; NMT 2015 from U. Montréal; NMT 2019 FAIR on newstest2019]

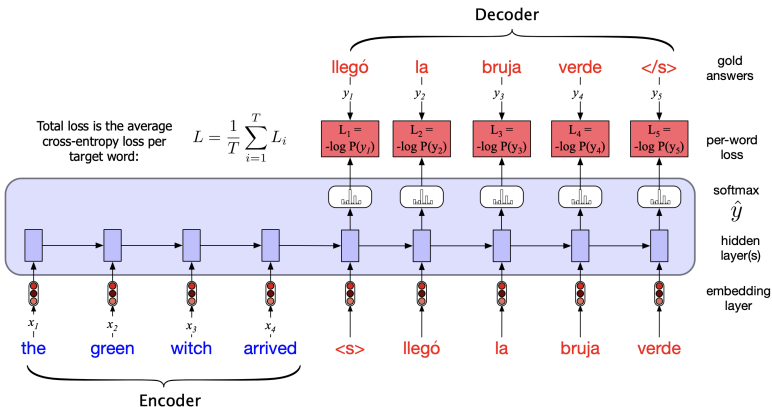


Sources: http://www.meta-net.eu/events/meta-forum-2016/slides/09_sennrich.pdf & <http://matrix.statmt.org/>

13

Training Seq2Seq

- Use **teacher-forcing** and compute the **average loss** over the predicted sequence
- Encoder-decoder architectures are trained **end-to-end**

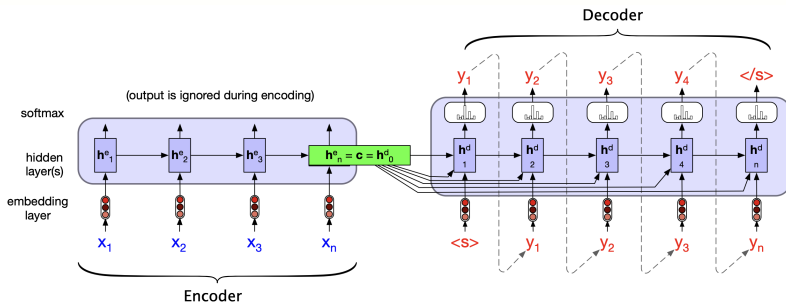


Propagating the context

- Problem:** the influence of the context, \mathbf{c} , decreases as the output sentence is generated
- Solution:** Make \mathbf{c} available at each step of the decoding process:


$$\mathbf{h}_t^d = g(\hat{\mathbf{y}}_{t-1}, \mathbf{h}_{t-1}^d, \mathbf{c}), \quad \text{for } t = 1, \dots, m \quad (4)$$

where $\mathbf{c} = f(\mathbf{h}_{1:n}^e)$ and $\mathbf{h}_0^d = \mathbf{c}$



Weaknesses of Seq2Seq

- **Problem:** the context \mathbf{c} from the encoder must represent everything about the meaning of the source text
 - Information at different parts of the source text may not be equally represented in the context vector
- The **attention mechanism** is a solution: allows the decoder to access **all** hidden states in the encoder, not just \mathbf{c}



References

- Gers, F. A., Schmidhuber, J., and Cummins, F. (2000). Learning to forget: Continual prediction with LSTM. *Neural Computation*, 12(10):2451–2471.
- Hochreiter, S. and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780.
- Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the National Academy of Sciences*, 79(8):2554–2558.
- Jurafsky, D. and Martin, J. H. (2019). Speech and language processing (3rd (draft) ed.).
- Manning, C., Socher, R., Fang, G. G., and Mundra, R. (2017). CS224n: Natural Language Processing with Deep Learning.
- Olah, C. (2015). Understanding LSTM Networks.
- Rumelhart, D. E., Hinton, G. E., and Williams, R. J. (1985). Learning internal representations by error propagation. Technical report, California Univ San Diego La Jolla Inst for Cognitive Science.